

```

In [3]: import tensorflow as tf
        from tensorflow import keras
        from keras.models import Sequential
        from keras.layers import Activation, Dense, Flatten, BatchNormalization, Conv2D, MaxPool2D, Dropout
        from keras.optimizers import Adam, SGD
        from keras.metrics import categorical_crossentropy
        from keras.preprocessing.image import ImageDataGenerator

        import warnings
        import numpy as np
        import cv2
        from keras.callbacks import ReduceLROnPlateau
        from keras.callbacks import ModelCheckpoint, EarlyStopping
        warnings.simplefilter(action='ignore', category=FutureWarning)

        background = None
        accumulated_weight = 0.5

        #Creating the dimensions for the ROI...
        ROI_top = 100
        ROI_bottom = 300
        ROI_right = 150
        ROI_left = 350

```

```

In [4]: train_path = r'/home/anonimouz/Documents/SignLanguage/gesture/train'
        test_path = r'/home/anonimouz/Documents/SignLanguage/gesture/test'
        train_batches = ImageDataGenerator(preprocessing_function=tf.keras.applications.vgg16.preprocess_input).flow_from_dir(train_path)
        test_batches = ImageDataGenerator(preprocessing_function=tf.keras.applications.vgg16.preprocess_input).flow_from_dir(test_path)

Found 31101 images belonging to 26 classes.
Found 730 images belonging to 26 classes.

```

```

In [17]: import matplotlib.pyplot as plt
         imgs, labels = next(train_batches)
         #Plotting the images...
         def plotImages(images_arr):
             fig, axes = plt.subplots(1, 10, figsize=(30,20))
             axes = axes.flatten()
             for img, ax in zip( images_arr, axes):
                 img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
                 ax.imshow(img)
                 ax.axis('off')
             plt.tight_layout()
             plt.show()
         plotImages(imgs)
         print(imgs.shape)
         print(labels)

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

```

```
print(labels)
```

[illegible]

(10, 64, 64, 3)

[illegible]

```
In [6]: model = Sequential()
model.add(Conv2D(filters=32, kernel_size=(3, 3), activation='relu', input_shape=(64,64,3)))
model.add(MaxPool2D(pool_size=(2, 2), strides=2))
model.add(Conv2D(filters=64, kernel_size=(3, 3), activation='relu', padding = 'same'))
model.add(MaxPool2D(pool_size=(2, 2), strides=2))
model.add(Conv2D(filters=128, kernel_size=(3, 3), activation='relu', padding = 'valid'))
model.add(MaxPool2D(pool_size=(2, 2), strides=2))
model.add(Flatten())
model.add(Dense(64,activation = "relu"))
model.add(Dense(128,activation = "relu"))
#model.add(Dropout(0.2))
model.add(Dense(128,activation = "relu"))
#model.add(Dropout(0.3))
model.add(Dense(26,activation = "softmax"))
```

```
In [7]: model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 62, 62, 32)	896
max_pooling2d_3 (MaxPooling2D)	(None, 31, 31, 32)	0
conv2d_4 (Conv2D)	(None, 31, 31, 64)	18496
max_pooling2d_4 (MaxPooling2D)	(None, 15, 15, 64)	0
conv2d_5 (Conv2D)	(None, 13, 13, 128)	73856
max_pooling2d_5 (MaxPooling2D)	(None, 6, 6, 128)	0
flatten_1 (Flatten)	(None, 4608)	0
dense_4 (Dense)	(None, 64)	294976
dense_5 (Dense)	(None, 128)	8320
dense_6 (Dense)	(None, 128)	16512
dense_7 (Dense)	(None, 26)	3354

=====
Total params: 416,410

Trainable params: 416,410

Non-trainable params: 0
=====

```
In [8]:
```

```
model.compile(optimizer=SGD(learning_rate=0.001), loss='categorical_crossentropy', metrics=['accuracy'])
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.2, patience=1, min_lr=0.0005)
early_stop = EarlyStopping(monitor='val_loss', min_delta=0, patience=2, verbose=0, mode='auto')
```

```
model.compile(optimizer=Adam(learning_rate=0.001), loss='categorical_crossentropy', metrics=['accuracy'])
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.2, patience=1, min_lr=0.0001)
early_stop = EarlyStopping(monitor='val_loss', min_delta=0, patience=2, verbose=0, mode='auto')
```

```
In [9]: history2 = model.fit(train_batches, epochs=10, callbacks=[reduce_lr, early_stop], validation_data = test_batches)
```

Epoch 1/10

3111/3111 [=====] - 76s 24ms/step - loss: 0.1676 - accuracy: 0.9638 - val_loss: 0.0613 - val_accuracy: 0.9863 - lr: 0.0010

Epoch 2/10

3111/3111 [=====] - 73s 24ms/step - loss: 0.0426 - accuracy: 0.9929 - val_loss: 1.6384e-05 - val_accuracy: 1.0000 - lr: 0.0010

Non-trainable params: 0

In [8]:

```
model.compile(optimizer=SGD(learning_rate=0.001), loss='categorical_crossentropy', metrics=['accuracy'])
reduce_lr = ReduceLRonPlateau(monitor='val_loss', factor=0.2, patience=1, min_lr=0.0005)
early_stop = EarlyStopping(monitor='val_loss', min_delta=0, patience=2, verbose=0, mode='auto')

model.compile(optimizer=Adam(learning_rate=0.001), loss='categorical_crossentropy', metrics=['accuracy'])
reduce_lr = ReduceLRonPlateau(monitor='val_loss', factor=0.2, patience=1, min_lr=0.0001)
early_stop = EarlyStopping(monitor='val_loss', min_delta=0, patience=2, verbose=0, mode='auto')
```

In [9]:

```
history2 = model.fit(train_batches, epochs=10, callbacks=[reduce_lr, early_stop], validation_data = test_batches)

Epoch 1/10
3111/3111 [=====] - 76s 24ms/step - loss: 0.1676 - accuracy: 0.9638 - val_loss: 0.0613 - v
al accuracy: 0.9863 - lr: 0.0010
Epoch 2/10
3111/3111 [=====] - 73s 24ms/step - loss: 0.0426 - accuracy: 0.9929 - val_loss: 1.6384e-05
- val_accuracy: 1.0000 - lr: 0.0010
Epoch 3/10
3111/3111 [=====] - 76s 24ms/step - loss: 0.0218 - accuracy: 0.9966 - val_loss: 1.1261e-06
- val_accuracy: 1.0000 - lr: 0.0010
Epoch 4/10
3111/3111 [=====] - 75s 24ms/step - loss: 6.0457e-05 - accuracy: 1.0000 - val_loss: 0.0061
- val_accuracy: 0.9959 - lr: 2.0000e-04
Epoch 5/10
3111/3111 [=====] - 73s 23ms/step - loss: 3.1256e-04 - accuracy: 0.9999 - val_loss: 1.2166
e-07 - val_accuracy: 1.0000 - lr: 1.0000e-04
Epoch 6/10
3111/3111 [=====] - 73s 23ms/step - loss: 3.4253e-07 - accuracy: 1.0000 - val_loss: 4.2295
e-08 - val_accuracy: 1.0000 - lr: 1.0000e-04
Epoch 7/10
3111/3111 [=====] - 74s 24ms/step - loss: 1.0593e-07 - accuracy: 1.0000 - val_loss: 9.9613
e-09 - val_accuracy: 1.0000 - lr: 1.0000e-04
Epoch 8/10
3111/3111 [=====] - 72s 23ms/step - loss: 2.5110e-08 - accuracy: 1.0000 - val_loss: 9.7980
e-10 - val_accuracy: 1.0000 - lr: 1.0000e-04
Epoch 9/10
3111/3111 [=====] - 73s 23ms/step - loss: 6.1941e-09 - accuracy: 1.0000 - val_loss: 0.0000
e+00 - val_accuracy: 1.0000 - lr: 1.0000e-04
Epoch 10/10
3111/3111 [=====] - 75s 24ms/step - loss: 1.7900e-09 - accuracy: 1.0000 - val_loss: 0.0000
e+00 - val_accuracy: 1.0000 - lr: 1.0000e-04
```

In [10]:

```
imgs, labels = next(test_batches)
```

In [11]:

```
scores = model.evaluate(imgs, labels, verbose=0)
print(f'{model.metrics_names[0]} of {scores[0]}; {model.metrics_names[1]} of {scores[1]*100}%')

loss of 0.0; accuracy of 100.0%
```

In [12]:

```
model.save('ASLmodel.h5')
```

```
In [10]: imgs, labels = next(test_batches)
```

```
In [11]: scores = model.evaluate(imgs, labels, verbose=0)
print(f'{model.metrics_names[0]} of {scores[0]}; {model.metrics_names[1]} of {scores[1]*100}%')
loss of 0.0; accuracy of 100.0%
```

```
In [12]: model.save('ASLmodel.h5')
```